KIEU STUDIO · GITHUB RADAR · EPISODE 5 · CHEAT SHEET

VoiceStudio

An AI voice studio that runs on your own device.

`github.com/debpalash/VoiceStudio`

+49,116	95,975	AGPL-3.0
stars in 30 days · ≈ 1637 stars/day	total stars · GitHub API	license · created Apr 9, 2026

Open-source app: clone a voice from a clean recording, design a voice from a description, dub video, transcribe, narrate stories and books, with a local API and MCP. The author says no paid cloud stands behind it.

How it works in six steps

1. **Install the app** — desktop, local processing
2. **Model Manager** — pick an engine, download only what you need
3. **Choose a voice** — clone from a clean recording, or describe one
4. **Text or video** — read aloud, dub, transcribe
5. **Runs on-device** — GPU if you have one, CPU works (slower)
6. **Automate** — local API, MCP, recipes, watch folders

Quick start

```
# macOS / Linux
curl -fsSL https://voicestudio.sh/install | sh

# Windows (PowerShell)
irm https://voicestudio.sh/install.ps1 | iex

# Skill for your agent
npx skills add debpalash/VoiceStudio

# After installing
Open Voice cloning → pick a voice or add a clean record
ing → type text → Generate

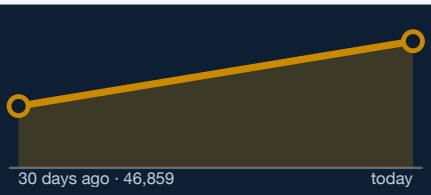
# Needs: A GPU is optional; models download on demand (size varies by model)
```

Example prompt: Install the VoiceStudio app on this device and verify it works, following https://github.com/debpalash/VoiceStudio/blob/main/docs/install/agent.md

Key commands / features

- **`Voice cloning`** — a familiar voice from one clean recording
- **`Voice design`** — describe a voice in words, then generate it
- **`Dubbing`** — translate and dub video, timed to the original
- **`Dictate & transcribe`** — floating dictation widget and local transcription
- **`Narrate & batch`** — stories, audiobooks, queued jobs

Growth



Two real points: the start of the window and today, axis from zero. The gain is as reported by GitHub Trending; we did not rebuild the star history.

Why this repo exists

- High-quality AI voices usually sit behind a cloud subscription: your text, and sometimes your voice, goes to someone else's server.
- VoiceStudio runs on your machine; the app asks before downloading any model.

What to use it for

- Narration for videos and podcasts on your own machine
- Dub a video into another language
- Transcribe locally, without uploading files
- Let an agent call voice features through the local API or MCP

Think twice before you use it

- Clone a voice only with its owner's permission (README).
- The app is AGPL-3.0; model weights, voices and generated output have their own licenses (see the Model Manager).
- The curl | sh installer runs code from the web: read the script before you run it.
- "646 languages" is the authors' own claim; Kieu Studio has not verified it.

Tried it? Kieu Studio uses a local VoiceStudio install (v0.5.6) for this series' narration; we have not tested every feature.

Sources: the repo's README, https://github.com/debpalash/VoiceStudio, https://voicestudio.sh. Stars: GitHub API; gains: GitHub Trending (as reported by Trending), snapshot Oct 31, 2026; stars measure attention, not quality: re-check before relying on them. Summarized in Kieu Studio's own words; commands copied from the README; no third-party logos or screenshots.
Disclosure: drafted with AI assistance; the voice and presenter in the videos are AI-generated. Not investment advice, and not endorsed by Kieu Studio. **kieu.app**